

Credit Default Prediction ¶

Dataset

Data Set Information: The training data contains 22500 observations with the predictor variables as well as the response variable. The test set contains 7500 observations with the response variable removed.

Task: Predict the response variable (default status) for the test data.

Variable descriptions: This research employed a binary variable, default payment (Yes = 1, No = 0), as the response variable.

This study reviewed the literature and used the following 23 variables as explanatory variables:

- **X1:** Amount of the given credit (NT dollar): it includes both the individual consumer credit and his/her family (supplementary) credit.
- **X2:** Gender (1 = male; 2 = female).
- **X3:** Education (1 = graduate school; 2 = university; 3 = high school; 4 = others).
- **X4:** Marital status (1 = married; 2 = single; 3 = others).
- **X5:** Age (year).
- **X6 - X11:** History of past payment. We tracked the past monthly payment records (from April to September, 2005) as follows:
 - **X6** = the repayment status in September, 2005;
 - **X7** = the repayment status in August, 2005;
 - **X11** = the repayment status in April, 2005. The measurement scale for the repayment status is:
 - -1 = pay duly;
 - 1 = payment delay for one month;
 - 2 = payment delay for two months;
 - 8 = payment delay for eight months;
 - 9 = payment delay for nine months and above.
 - -2 = indicates no consumption in the month, and a value of
 - 0 = indicates the use of revolving credit (equivalent to prepayment)
- **X12-X17:** Amount of bill statement (NT dollar).
 - **X12** = amount of bill statement in September, 2005;
 - **X13** = amount of bill statement in August, 2005;
 - **X17** = amount of bill statement in April, 2005.
- **X18-X23:** Amount of previous payment (NT dollar).
 - **X18** = amount paid in September, 2005;
 - **X19** = amount paid in August, 2005;
 - **X23** = amount paid in April, 2005.

Step 1: Data Import

In [43]: %%time

```
# Import the required Python Packages
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import warnings
from io import StringIO

# Import the AWS & Sagemaker Packages
import boto3
import sagemaker
from sagemaker import get_execution_role
from sagemaker.amazon.amazon_estimator import get_image_uri
from sagemaker.predictor import csv_serializer, json_deserializer

warnings.filterwarnings('ignore') # to suppress seaborn warnings
pd.options.display.max_columns = None
```

CPU times: user 76 µs, sys: 6 µs, total: 82 µs
Wall time: 86.8 µs

a) Get Role, Region, Session

```
In [30]: role = get_execution_role()
region = boto3.Session().region_name
sess = sagemaker.Session()

print("Role : ", role)
print("Region : ", region)
```

Role : arn:aws:iam::789247493478:role/SageMakerFullAccess
Region : us-east-1

b) Define S3 Bucket & store raw files

```
In [51]: bucket='g-demo' # put your s3 bucket name here, and create s3 bucket
        prefix = 'sagemaker'

        # customize to your bucket where you have stored the data
        bucket_path = 'https://s3-{}.amazonaws.com/{}'.format(region,bucket)
        s3= boto3.Session(region_name=region).resource('s3')

        def upload_to_s3(fobj,channel,filename):
            key = prefix+'/' +channel
            url = 's3://{}/{}/{}'.format(bucket, key, filename)
            print('Writing to {}'.format(url))
            s3.Bucket(bucket).Object(key).put(Body=fobj)

        upload_to_s3(open('credit_card_default_TRAIN.csv', 'rb'),'input/train','credit_card')
        upload_to_s3(open('credit_card_default_TEST.csv', 'rb'),'input/test','credit_card')
```

Writing to s3://g-demo/sagemaker/input/train/credit_card_default_TRAIN.csv
 Writing to s3://g-demo/sagemaker/input/test/credit_card_default_TEST.csv

Step 2: Data Preprocessing

a) Read the data and normalize it

```
In [52]: # Read the train & test data into DataFrames.
        train = pd.read_csv("credit_card_default_TRAIN.csv",index_col=0)
        test = pd.read_csv("credit_card_default_TEST.csv",index_col=0)

        # Fix Header of the data, row 0 serves as more sensible header names
        def fix_header(data):
            new_header = data.iloc[0] # take the first row for the header
            data = data[1:] # take the data without the header row
            data.columns = new_header # set the header row as the df header
            data.rename(columns={'default payment next month':'DEFAULTER'}, inplace=True)
            return data

        train = fix_header(train)
        test = fix_header(test)
```

Check for null values in the datasets.

```
In [53]: train.isnull().values.any(),test.isnull().values.any()
```

```
Out[53]: (False, False)
```

Both train and test datasets do not have null values

In [54]: `train.describe()`

Out[54]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5
count		22500	22500	22500	22500	22500	22500	22500	22500	22500	22500
unique		77	2	7	4	55	11	11	11	11	10
top		50000	2	2	2	29	0	0	0	0	0
freq		2630	13729	10634	12219	1249	11069	11849	11876	12556	12759

From above description of train data, we found that

- EDUCATION has 7 unique values instead of 4 (1 = graduate school; 2 = university; 3 = high school; 4 = others)
- MARRIAGE has 4 unique values instead of 3 (1 = married; 2 = single; 3 = others)

In [36]: `train.EDUCATION.value_counts()`

Out[36]:

```
2    10634
1     7982
3     3581
5      184
4       76
6        33
0        10
Name: EDUCATION, dtype: int64
```

According to description we should have values 1,2,3,4 thus we will change 5,6,0 to 4 i.e. others

```
In [55]: train.EDUCATION[train.EDUCATION=='0']='4'
train.EDUCATION[train.EDUCATION=='5']='4'
train.EDUCATION[train.EDUCATION=='6']='4'
train.EDUCATION.unique()

test.EDUCATION[test.EDUCATION=='0']='4'
test.EDUCATION[test.EDUCATION=='5']='4'
test.EDUCATION[test.EDUCATION=='6']='4'
test.EDUCATION.unique()
```

Out[55]: `array(['2', '1', '3', '4'], dtype=object)`

```
In [38]: train.MARRIAGE.value_counts()
```

```
Out[38]: 2    12219
         1     9990
         3      255
         0       36
         Name: MARRIAGE, dtype: int64
```

```
In [56]: # According to description we should have values 1,2,3 thus we will change 0 to 3
         train.MARRIAGE[train.MARRIAGE=='0']='3'
         train.MARRIAGE.unique()

         test.MARRIAGE[test.MARRIAGE=='0']='3'
         test.MARRIAGE.unique()
```

```
Out[56]: array(['1', '2', '3'], dtype=object)
```

Change target variable(DEFAULTER) data-type as "int" and put it at first position

```
In [57]: train.DEFAULTER = train.DEFAULTER.astype(int)
         cols = list(train.columns)
         cols = [cols[-1]] + cols[:-1]
         train = train[cols]
         train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 22500 entries, 1 to 22500
Data columns (total 24 columns):
DEFAULTER    22500 non-null int64
LIMIT_BAL   22500 non-null object
SEX          22500 non-null object
EDUCATION    22500 non-null object
MARRIAGE     22500 non-null object
AGE          22500 non-null object
PAY_0        22500 non-null object
PAY_2        22500 non-null object
PAY_3        22500 non-null object
PAY_4        22500 non-null object
PAY_5        22500 non-null object
PAY_6        22500 non-null object
BILL_AMT1    22500 non-null object
BILL_AMT2    22500 non-null object
BILL_AMT3    22500 non-null object
BILL_AMT4    22500 non-null object
BILL_AMT5    22500 non-null object
BILL_AMT6    22500 non-null object
PAY_AMT1     22500 non-null object
PAY_AMT2     22500 non-null object
PAY_AMT3     22500 non-null object
PAY_AMT4     22500 non-null object
PAY_AMT5     22500 non-null object
PAY_AMT6     22500 non-null object
dtypes: int64(1), object(23)
memory usage: 4.3+ MB
```

In [58]: `train.head()`

Out[58]:

	ID	DEFAULTER	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
Variable											
1		1	20000	2	2	1	24	2	2	-1	
2		1	120000	2	2	2	26	-1	2	0	
3		0	90000	2	2	2	34	0	0	0	
4		0	50000	2	2	1	37	0	0	0	
5		0	50000	1	2	1	57	-1	0	-1	

b) Upload the normalized data to S3

```
In [59]: train_str_buffer = StringIO()
train.to_csv(train_str_buffer, index=False,header=None)

test_str_buffer = StringIO()
test.to_csv(test_str_buffer, index=False,header=None)

upload_to_s3(train_str_buffer.getvalue(),'normalized/train','credit_card_default_TRAIN.csv')
upload_to_s3(test_str_buffer.getvalue(),'normalized/test','credit_card_default_TEST.csv')

Writing to s3://g-demo/sagemaker/normalized/train/credit_card_default_TRAIN.csv
Writing to s3://g-demo/sagemaker/normalized/test/credit_card_default_TEST.csv
```

Step 3: Data Visualization

Check correlation between different features using heatmap

```
In [60]: cor = train.astype(float).corr()
plt.show()
plt.figure(figsize=(18,18))
sns.heatmap(cor, cbar = True, square = True, annot=True, fmt= '.2f',annot_kws={
    xticklabels=cor.columns.values,
    yticklabels=cor.columns.values})
```



Above heatmap shows that

- 'BILL_AMTX' are highly correlated to each other, but very less correlation to target label 'DEFAULTER'. These can be removed from normalized data when large data is there.
- Payment statuses 'PAY' show highest contribution to the defaulter label.
- We can see above that PAY_0,PAY_2...have high positive correlation to DEFAULTER
- LIMIT_BAL has pretty high negative correlation**

Predictive Modelling (AWS Sagemaker)

Step 4: Get the Linear-Learner Algo container

```
In [85]: container = get_image_uri(boto3.Session().region_name, 'linear-learner', "latest")
print("Container : ",container)
```

Container : 382416733822.dkr.ecr.us-east-1.amazonaws.com/linear-learner:latest

Step 5 : Build the model

```
In [97]: s3_train_data = 's3://{}/{}/normalized/{}'.format(bucket, prefix, 'train')
print('Normalized training data location: {}'.format(s3_train_data))

output_location = 's3://{}/{}/{}'.format(bucket, prefix)
print('Training model artifact will be uploaded to: {}'.format(output_location))

linear_learner_model = sagemaker.estimator.Estimator(container,
                                                        role,
                                                        train_instance_count=1,
                                                        train_instance_type='ml.m4.xlarge',
                                                        base_job_name='credit-default-prediction',
                                                        output_path=output_location,
                                                        sagemaker_session=sess
                                                        )
```

Normalized training data location: s3://g-demo/sagemaker/normalized/train
 Training model artifact will be uploaded to: s3://g-demo/sagemaker

Hyperparameters tuning (Optional)

- feature_dim is set to 23 excluding target label = DEFAULTER
- predictor_type is set to 'binary_classifier' since we are trying to predict whether the defaultter is 1(Yes) or 0 (No).
- mini_batch_size is set to 1000. This value can be tuned for relatively minor improvements in fit and speed, but selecting a reasonable value relative to the dataset is appropriate in most cases.

```
In [99]: linear_learner_model.set_hyperparameters(feature_dim=23,
                                                    predictor_type='binary_classifier',
                                                    optimizer: "adam",
                                                    mini_batch_size=1000)
```

Step 6: Train the model


```
In [102]: # Create the training job to train the model
train_channel = sagemaker.session.s3_input(s3_train_data, content_type='text/csv')
linear_learner_model.fit({'train': train_channel}, logs=False, job_name='credit-de'
```

```
2019-09-14 20:54:37 Starting - Starting the training job
2019-09-14 20:54:41 Starting - Launching requested ML instances.....
2019-09-14 20:55:43 Starting - Preparing the instances for trainin
g.....
2019-09-14 20:56:54 Downloading - Downloading input data...
2019-09-14 20:57:17 Training - Downloading the training image..
2019-09-14 20:57:29 Training - Training image download completed. Training in p
rogress...
2019-09-14 20:57:45 Uploading - Uploading generated training model
2019-09-14 20:57:51 Completed - Training job completed
```

Step 7: Deploy the model for real time predictions

```
In [104]: # Real time Prediction
predictor = linear_learner_model.deploy(
    initial_instance_count=1,
    instance_type='ml.m4.xlarge',
    endpoint_name="credit-default-prediction"
)
```

Using already existing model: credit-default-prediction

```
-----
-----!
```

a) Realtime predictions : predict() method

```
In [135]: single_test = r'20000,2,2,1,24,2,2,-1,-1,-2,-2,3913,3102,689,0,0,0,0,689,0,0,0,0

predictor.content_type = 'text/csv'
predictor.serializer = csv_serializer
predictor.deserializer = json_deserializer

result = predictor.predict(single_test)
print('Input Record:: ',single_test)
print('Prediction :: ',result)
```

```
Input Record:: 20000,2,2,1,24,2,2,-1,-1,-2,-2,3913,3102,689,0,0,0,0,689,0,0,0,0
0
Prediction :: {'predictions': [{'score': 0.4804583191871643, 'predicted_label': 1.0}]}
```

b) Real time predictions : Invoking endpoint by name

```
In [138]: runtime_client = boto3.client('runtime.sagemaker')
          endpoint_name = 'credit-default-prediction'

          response = runtime_client.invoke_endpoint(EndpointName = endpoint_name,
                                                    ContentType = 'text/csv',
                                                    Body = single_test)
          result = response['Body'].read().decode('utf-8')
          print('Input Record:: ',single_test)
          print('Prediction :: ',result)
```

```
Input Record::  20000,2,2,1,24,2,2,-1,-1,-2,-2,3913,3102,689,0,0,0,0,689,0,0,0,
0
Prediction ::  {"predictions": [{"score": 0.4804583191871643, "predicted_label": 1.0}]}
```

Step 8: Batch Transform Predictions

```
In [144]: # Create a batch transform job
          batch_input = 's3://{}/{}/normalized/{}'.format(bucket, prefix, 'test') # The Location of the Input Data
          batch_output = 's3://{}/{}/{}/{}'.format(bucket, prefix, 'output') # The Location of the Output Data

          batch_transformer = linear_learner_model.transformer(instance_count=1,
                                                                instance_type='ml.m4.xlarge',
                                                                output_path=batch_output,accept='Line',
                                                                assemble_with='Line')

          batch_transformer.transform(data=batch_input,
                                     data_type='S3Prefix',
                                     content_type='text/csv',
                                     split_type='Line',
                                     join_source='Input'
                                     #,job_name='batch-credit-default-predictions'
                                     )
          batch_transformer.wait(logs=False)
          print("Batch predictions generated successfully !!")
```

Using already existing model: credit-default-prediction

```
.....!
Batch predictions generated successfully !!
```

Model Metrics

```
In [167]: metrics_dataframe = linear_learner_model.training_job_analytics.dataframe()  
metrics_dataframe[['metric_name', 'value']]
```

```
WARNING:root:Warning: No metrics called test:binary_f_beta found  
WARNING:root:Warning: No metrics called test:mse found  
WARNING:root:Warning: No metrics called validation:binary_f_beta found  
WARNING:root:Warning: No metrics called validation:objective_loss found  
WARNING:root:Warning: No metrics called validation:objective_loss:final found  
WARNING:root:Warning: No metrics called test:absolute_loss found  
WARNING:root:Warning: No metrics called train:mse found  
WARNING:root:Warning: No metrics called validation:recall found  
WARNING:root:Warning: No metrics called validation:precision found  
WARNING:root:Warning: No metrics called test:recall found  
WARNING:root:Warning: No metrics called test:objective_loss found  
WARNING:root:Warning: No metrics called test:precision found  
WARNING:root:Warning: No metrics called validation:mse found  
WARNING:root:Warning: No metrics called validation:binary_classification_accuracy found  
WARNING:root:Warning: No metrics called train:absolute_loss found  
WARNING:root:Warning: No metrics called test:binary_classification_accuracy found  
WARNING:root:Warning: No metrics called validation:absolute_loss found
```

Out[167]:

	metric_name	value
0	train:progress	45.428571
1	train:objective_loss	0.487678
2	train:recall	0.381682
3	train:precision	0.630110
4	train:objective_loss:final	0.472193
5	train:binary_f_beta	0.475398
6	train:binary_classification_accuracy	0.809511
7	train:throughput	29202.586217